

Hitchhiker's Guide to the Tidyverse

Cory Lanker

2019-07-22

Contents

Introduction	5
1 tibbles, ggplot2, and the <i>tidyverse</i>	7
1.1 Tibbles: Boston housing data	8
1.2 ggplot2 and EDA	11
1.3 Many plotting options	21
2 dplyr and tidyr	29
2.1 Hoofin' it with dplyr	29
2.2 tidyr and relational data	31

Introduction

Introducing the tidyverse with two data sets:

1. Basic plots with `tibble` and `ggplot2` using `Boston` house prices.
2. Preprocessing with `tidyr` and `dplyr` using `Lahman` baseball data.

Though some of these commands will be used, we won't go deeply into the following tidyverse packages:

1. Reading in data with `readr`.
2. String manipulation with `stringr`.
3. Dates and times with `lubridate`.
4. Handling factors with `forcats`.

R proficiency is assumed. These notes aim to bring a functional R coder into the tidyverse realm.

```
# To install the necessary packages in the tidyverse:  
install.packages("tidyverse", dependencies = TRUE)
```

This document is built with R Markdown, **knitr** (Xie, 2015), and the **bookdown** package (Xie, 2019).

Chapter 1

tibbles, ggplot2, and the *tidyverse*

The tidyverse universe includes:

In general, the tidyverse is the following:

1. provided the `pipe` command `%>%`
 - `x %>% f(y, z, ...)` is `f(x, y, z, ...)`
 - allows chained commands for better coherence
 - e.g., `mtcars %>% apply(2, mean)` is error without `tidyr::%>%`
2. `tibble` is the improved data structure of the tidyverse
 - easier to read-in data to a useful format
 - automatic type conversion
 - nicer printing options
3. `dplyr` provides tibble manipulation commands
 - understandable data processing with `pipe` streams
 - **filter** data faster
 - **arrange** rows of data easily
 - **select** columns quickly
 - **mutate** variables
 - **summarize** according to `group_by()`
 - also provides SQL relational operations
4. `ggplot2` is a plotting syntax (grammar of graphics)
 - `qplot()` provides a sensible **quick plot**
 - apply plot types to data rather than the reverse

- e.g. `ggplot(data) + plot_type(aes(xvar, yvar, groups), options)`
 - allows grid of plots by group using **facets**
 - overlays statistical summaries, e.g. `+ geom_smooth(x, y)`
 - “add” options such as transformed axes, labels, coordinates, etc.
5. **readr** is a faster, less painful read-in method
- **read_fun** denotes **readr** functions (instead of **read.fun**)
 - guesses column types
 - offers writing functions, too
 - allows read and write with RDS, R’s binary format
6. **tidyr** recharacterizes tibbles
- **spread()** turns key and value columns into key-category columns
 - e.g., **state**, **year**, **pop** into **state**, **1990**, **1991**, ... of **pop** values
 - **gather()** turns expands data frames by condensing columns
 - e.g., condenses **1990**, **1991**, ... into two **year**, **pop** columns
7. Other helpful tidyverse packages:
- **stringr** offers many useful **str_fun** operations
 - **forcats** has operations for categorical variables
 - **lubridate** provides date and time control
 - **purrr**

The examples I’ll use is the Boston housing database and the Lahman baseball database. By doing analysis on these two data sets, I hope to introduce the power of the tidyverse.

1.1 Tibbles: Boston housing data

Load, convert, print a tibble.

```
# Convert to a tibble so it prints nicely
library(MASS)
select <- dplyr::select
boston <- as_tibble(MASS::Boston)
boston
```

```
## # A tibble: 506 x 14
##      crim    zn indus  chas   nox    rm   age   dis   rad   tax ptratio
##      <dbl> <dbl> <dbl> <int> <dbl> <dbl> <dbl> <dbl> <int> <dbl>    <dbl>
##  1 0.00632  18    2.31     0 0.538  6.58  65.2  4.09     1   296    15.3
##  2 0.0273   0    7.07     0 0.469  6.42  78.9  4.97     2   242    17.8
##  3 0.0273   0    7.07     0 0.469  7.18  61.1  4.97     2   242    17.8
##  4 0.0324   0    2.18     0 0.458  7.00  45.8  6.06     3   222    18.7
```

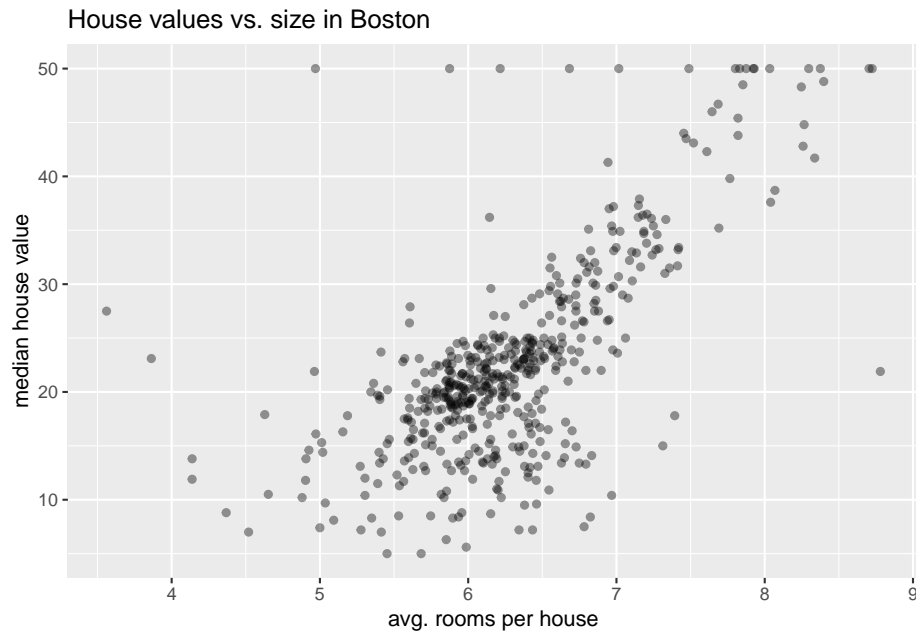


```
## 5 0.0690    0    2.18    0 0.458  7.15  54.2  6.06    3  222   18.7
## 6 0.0298    0    2.18    0 0.458  6.43  58.7  6.06    3  222   18.7
## 7 0.0883   12.5  7.87    0 0.524  6.01  66.6  5.56    5  311   15.2
## 8 0.145    12.5  7.87    0 0.524  6.17  96.1  5.95    5  311   15.2
## 9 0.211    12.5  7.87    0 0.524  5.63 100    6.08    5  311   15.2
## 10 0.170   12.5  7.87    0 0.524  6.00  85.9  6.59    5  311   15.2
## # ... with 496 more rows, and 3 more variables: black <dbl>, lstat <dbl>,
## #   medv <dbl>
?MASS::Boston
```

- crim per capita crime rate by town.
- zn proportion of residential land zoned for lots over 25,000 sq.ft.
- indus proportion of non-retail business acres per town.
- chas Charles River dummy variable (= 1 if tract bounds river; 0 otherwise).
- nox nitrogen oxides concentration (parts per 10 million).
- rm average number of rooms per dwelling.
- age proportion of owner-occupied units built prior to 1940.
- dis weighted mean of distances to five Boston employment centres.
- rad index of accessibility to radial highways.
- tax full-value property-tax rate per \$10,000.
- ptratio pupil-teacher ratio by town.
- black $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town.
- lstat lower status of the population (percent).
- medv median value of owner-occupied homes in \$1000s.

A ggplot is the first declaration (usually variable `data` is defined), followed by graphics definitions (operations on the data):

```
ggplot(data = boston) +
  geom_point(mapping = aes(x = rm, y = medv), alpha=0.4) +
  labs(x = "avg. rooms per house",
       y = "median house value",
       title = "House values vs. size in Boston")
```



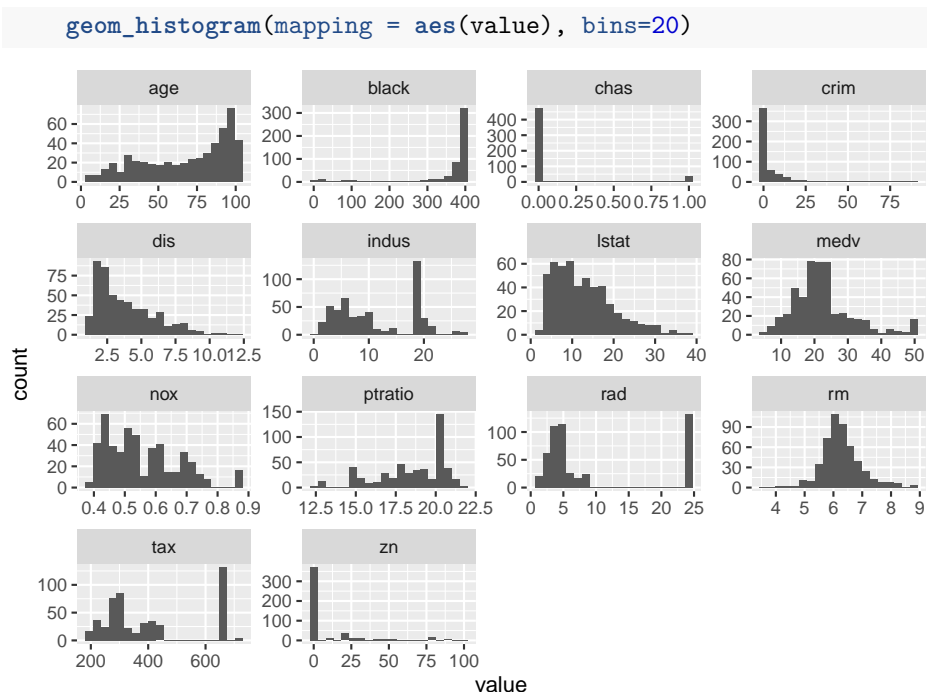
Making a histogram of all numeric variables. First step, gather all variables.

```
boston %>%
  keep(is.numeric) %>% # strips all non-numeric columns (unnecessary here)
  gather() # puts all variable values in a single column 'value'
```

```
## # A tibble: 7,084 x 2
##   key      value
##   <chr>   <dbl>
## 1 crim  0.00632
## 2 crim  0.0273
## 3 crim  0.0273
## 4 crim  0.0324
## 5 crim  0.0690
## 6 crim  0.0298
## 7 crim  0.0883
## 8 crim  0.145
## 9 crim  0.211
## 10 crim 0.170
## # ... with 7,074 more rows
```

Facet wrap allows plotting each key level separately.

```
boston %>%
  gather() %>%
  ggplot() +
    facet_wrap(~ key, scales = "free") +
```



From the histograms, there seems to be only a few values of `crim` over 30.

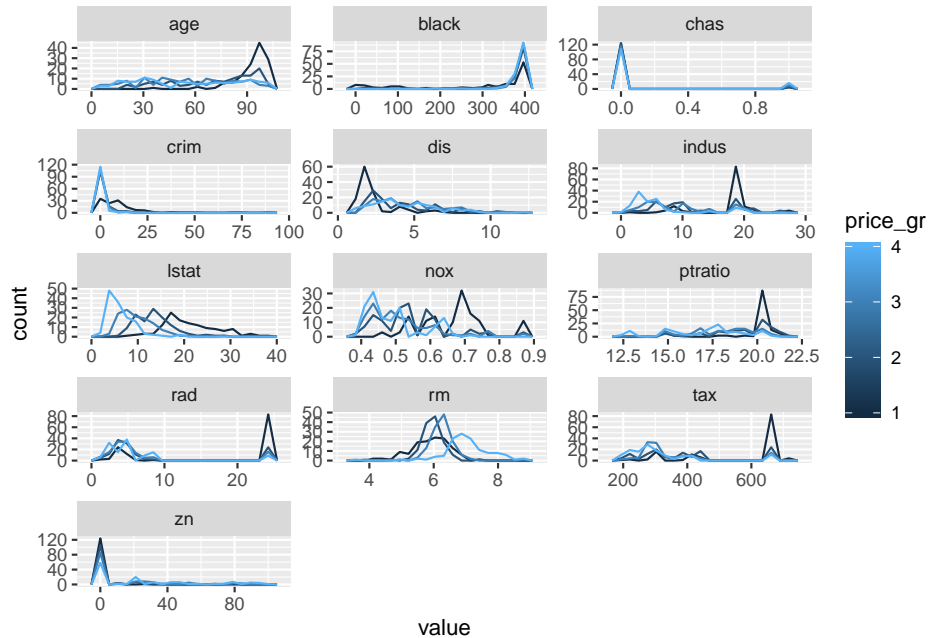
```
boston %>%  
  filter(crim > 30)
```

```
## # A tibble: 8 x 14  
##   crim    zn indus  chas  nox    rm  age  dis  rad  tax ptratio black  
##   <dbl> <dbl> <dbl> <int> <dbl> <dbl> <dbl> <dbl> <int> <dbl> <dbl> <dbl>  
## 1  89.0    0  18.1    0 0.671  6.97  91.9  1.42   24  666   20.2  397.  
## 2  38.4    0  18.1    0 0.693  5.45  100   1.49   24  666   20.2  397.  
## 3  41.5    0  18.1    0 0.693  5.53  85.4  1.61   24  666   20.2  329.  
## 4  67.9    0  18.1    0 0.693  5.68  100   1.43   24  666   20.2  385.  
## 5  51.1    0  18.1    0 0.597  5.76  100   1.41   24  666   20.2   2.6  
## 6  45.7    0  18.1    0 0.693  4.52  100   1.66   24  666   20.2  88.3  
## 7  73.5    0  18.1    0 0.679  5.96  100   1.80   24  666   20.2  16.4  
## 8  37.7    0  18.1    0 0.679  6.20  78.7  1.86   24  666   20.2  18.8  
## # ... with 2 more variables: lstat <dbl>, medv <dbl>
```

1.2 ggplot2 and EDA

But we want to know the conditional distributions according to `medv`. First, showing this with conditional densities.

```
boston %>%
  gather('key', 'value', -one_of('medv')) %>%
  mutate(price_gr = ntile(medv, 4)) %>%
  ggplot(aes(value, color = price_gr, group = price_gr)) +
    facet_wrap(~ key, ncol = 3, scales = "free") +
    geom_freqpoly(bins = 20)
```



Click on the expand icon at the top right to make bigger.

Appears chas is categorical.

```
boston <- boston %>%
  mutate(chas = factor(chas))
```

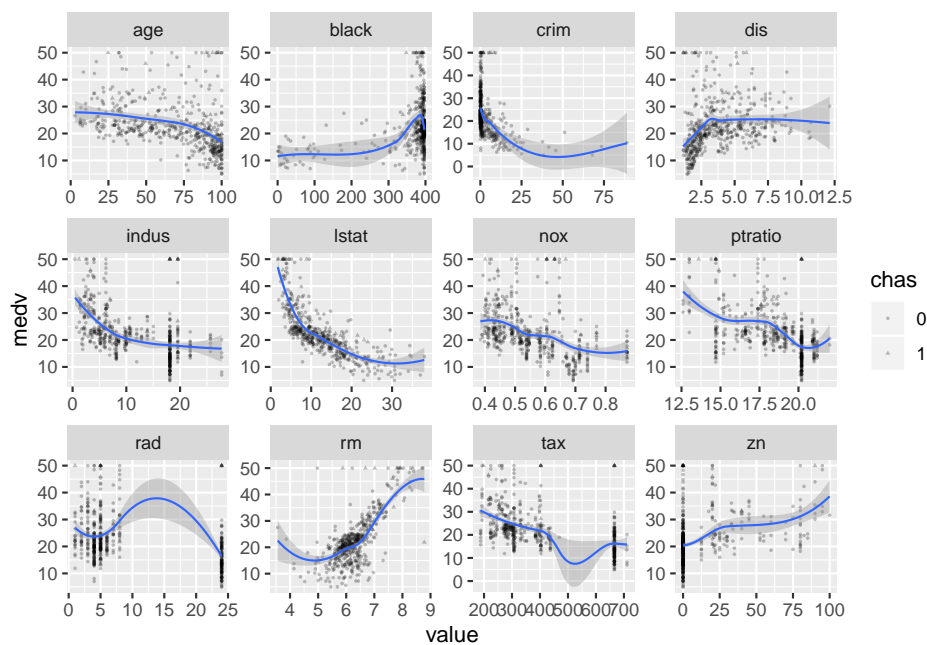
Second, scatterplots of median value vs. all variables.

```
boston %>%
  gather('key', 'value', -one_of(c("medv", "chas"))) %>%
  ggplot(aes(x = value, y = medv)) +
    facet_wrap(~ key, scales = "free") +
    geom_point(aes(shape = chas), size = 0.5, alpha = 0.25) +
    geom_smooth(lwd = 0.5, se = TRUE) +
    ggsave('plots/medv-scatter.pdf')
```

Saving 6.5 x 4.5 in image

`geom_smooth()` using method = 'loess' and formula 'y ~ x'

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

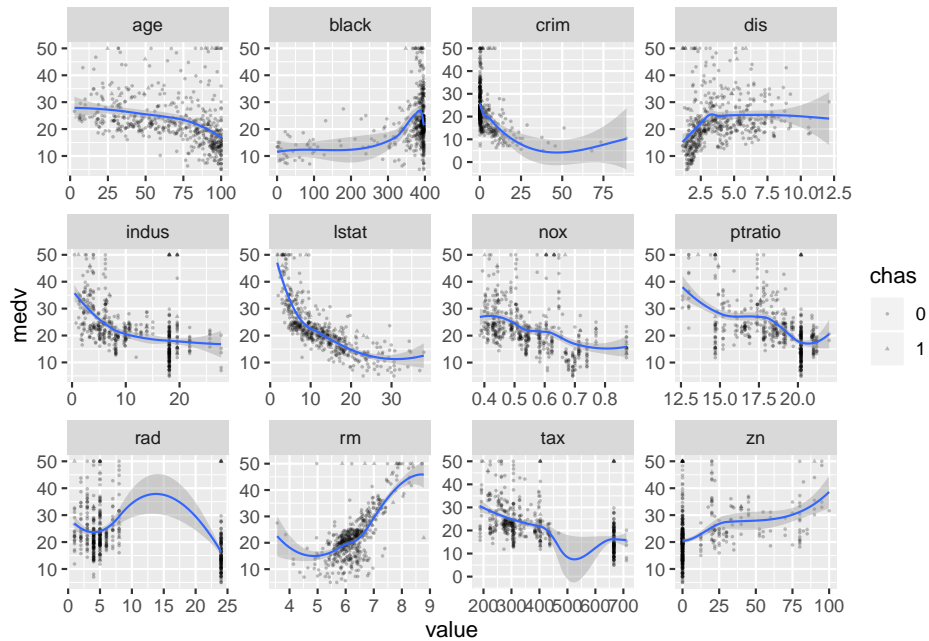


Click on the expand icon at the top right to make bigger.

There are ggplot jitter options, but none worked for me.

```
boston %>%
  gather('key', 'value', -one_of(c("medv", "chas"))) %>%
  ggplot(aes(x = value, y = medv)) +
    facet_wrap(~ key, scales = "free") +
    geom_jitter(aes(shape = chas), size = 0.5, alpha = 0.25) +
    geom_smooth(lwd = 0.5, se = TRUE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

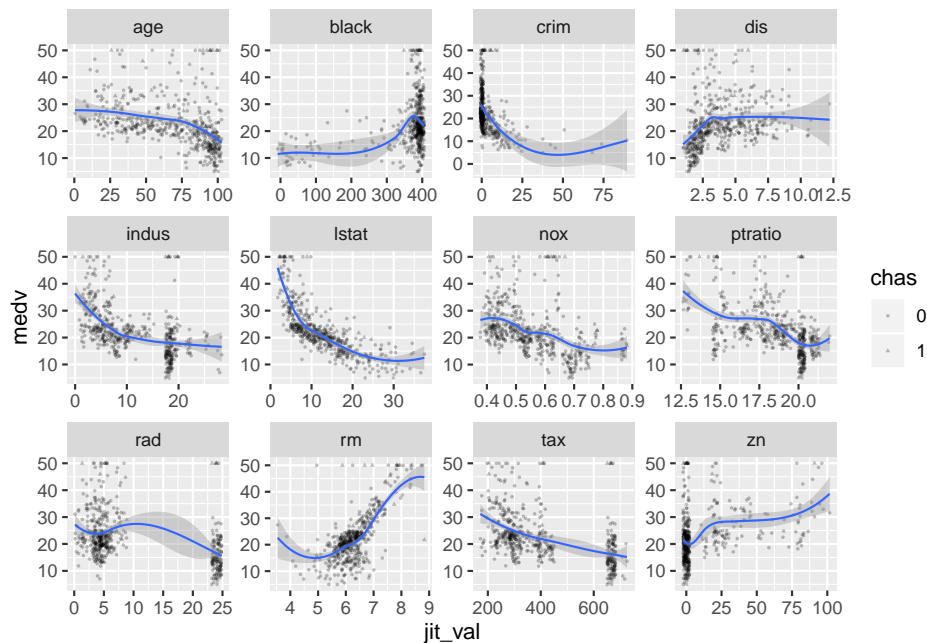


Tinkering to get a jittered plot.

```
var_sd <- boston %>%
  gather('key', 'value', -one_of(c("medv", "chas"))) %>%
  group_by(key) %>%
  summarize(var_sd = sd(value))
#var_sd <- boston %>%
# keep(is.numeric) %>%
# summarize_all(sd)
boston %>%
  gather('key', 'value', -one_of(c("medv", "chas"))) %>%
  left_join(y = var_sd, by = "key") %>%
  mutate(jit_val = value + var_sd * runif(nrow(boston), -0.1, 0.1)) %>%
  ggplot(aes(x = jit_val, y = medv)) +
    facet_wrap(~ key, scales = "free") +
    geom_jitter(aes(shape = chas), size = 0.5, alpha = 0.25) +
    geom_smooth(lwd = 0.5, se = TRUE) +
    ggsave('plots/medv-jitter.pdf')
```

```
## Saving 6.5 x 4.5 in image
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Click on the expand icon at the top right to make bigger.

Covariance plot of variables.


```
boston_num <- boston %>%
  keep(is.numeric)
boston_num %>%
  cor() %>%
  as_tibble() %>%
  mutate(name = colnames(boston_num)) %>%
  gather( , , -one_of("name")) %>%
  ggplot(aes(name, key, fill = value)) +
  geom_tile() +
  scale_fill_gradient2(low = "blue", mid = "white", high = "red",
    breaks = seq(-1, 1, by = 0.2)) +
  theme(legend.key.height = unit(45, "pt"))
```



Analyze median value and highway access `rad`.

```
boston %>%
  ggplot(aes(rad, medv)) +
  geom_jitter(aes(color = chas),
    height = 2, width = NULL,
    size = 1, alpha = 0.4) +
  geom_smooth(aes(color = chas), lwd = 1)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

Perhaps `rad = 24` is a missing value.

```
boston %>%
  count(rad)
```

```
## # A tibble: 9 x 2
##   rad     n
##   <int> <int>
## 1     1    20
## 2     2    24
## 3     3    38
## 4     4   110
## 5     5   115
## 6     6    26
## 7     7    17
## 8     8    24
## 9    24   132
```

```
boston %>%
  gather( , , -one_of("rad")) %>%
  group_by(key, rad) %>%
  mutate(value = as.numeric(value)) %>% # necessary due to factor variable chas
  summarize(z = round(mean(value), 1)) %>%
  spread(rad, z)
```

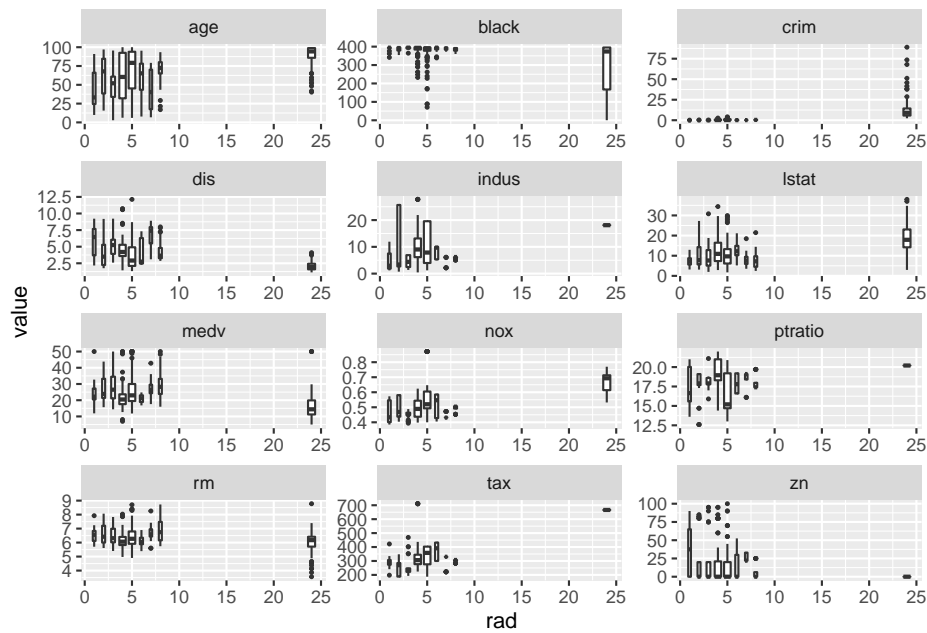
```
## Warning: attributes are not identical across measure variables;
## they will be dropped
```

```
## # A tibble: 13 x 10
## # Groups:   key [13]
##   key      `1`      `2`      `3`      `4`      `5`      `6`      `7`      `8`      `24`
##   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 age      45    64.8  49.3  60.8  69.2  60.1  40.1  67.3  89.8
## 2 black  389.   386.  392.  383.  369.  387.  388.  385.  288.
## 3 chas     0     0    0.1   0.1   0.1   0    0    0.2   0.1
## 4 crim     0    0.1   0.1   0.4   0.7   0.2   0.2   0.4  12.8
## 5 dis       6    4.1   5.1   4.4   3.7   4    6.5   4.4   2.1
## 6 indus     5.1   9.6   4.4  10.7   9.8   8.2   5    5.9  18.1
## 7 lstat     7.4  10    9.1  12.2  10.7  12.3   8     8   18.6
## 8 medv     24.4  26.8  27.9  21.4  25.7  21    27.1  30.4  16.4
## 9 nox      0.5   0.5   0.5   0.5   0.6   0.5   0.4   0.5   0.7
## 10 ptratio 17.6  17.3  18.2  19.1  16.5  17.8  18.4  18    20.2
## 11 rm       6.6   6.6   6.5   6.1   6.4   6.1   6.6   7     6
## 12 tax     291.  261.  246.  336   332.  373.  304.  301.  666
## 13 zn      39.9  20.4  16.4  14.7  11.1  13    26.7   6.2   0
```

Or in helpful boxplot format.

```
boston %>%
  keep(is.numeric) %>%
  gather( , , -one_of("rad")) %>%
  group_by(key, rad) %>%
  ggplot(aes(x = rad, y = value, group = rad)) +
    geom_boxplot(outlier.size = 0.5, varwidth = T) +
    facet_wrap(~ key, ncol = 3, scales = "free") +
    ggsave('plots/rad-boxplot.pdf')
```

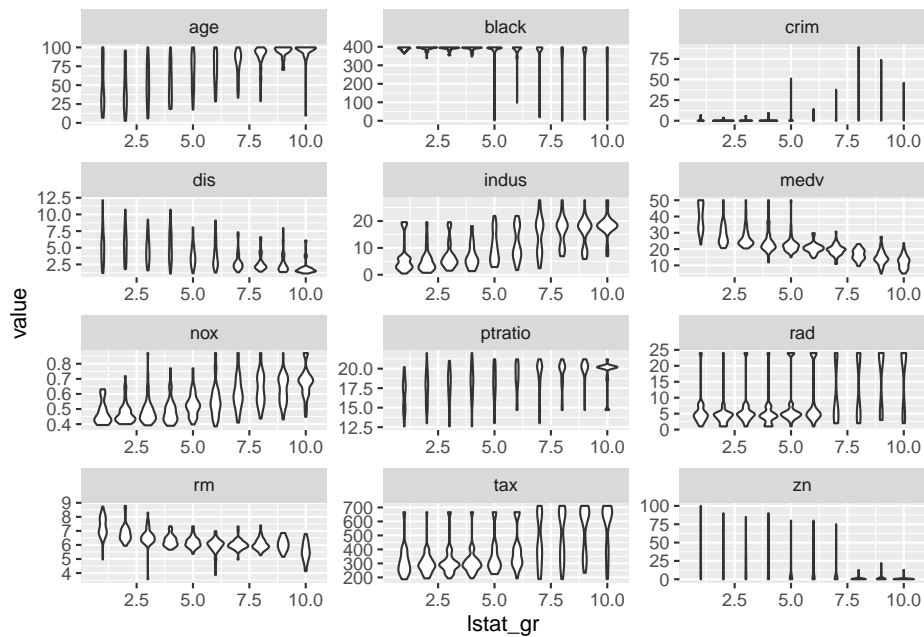
```
## Saving 6.5 x 4.5 in image
```



Looking at lstat relationships.


```
boston %>%
  keep(is.numeric) %>%
  gather( , , -one_of("lstat")) %>%
  mutate(lstat_gr = ntile(lstat, 10)) %>%
  group_by(key, lstat_gr) %>%
  ggplot(aes(x = lstat_gr, y = value, group = lstat_gr)) +
    geom_violin() +
    facet_wrap(~ key, ncol = 3, scales = "free") +
    ggsave('plots/lstat-violin.pdf')
```

Saving 6.5 x 4.5 in image



Jittering works well for single plots.

```
boston %>%
  ggplot(aes(tax, nox)) +
    geom_jitter(aes(color = medv, shape = chas),
               height = 0.02, width = 10) +
    scale_color_gradient2(midpoint = 20,
                          low = "blue", mid = "gray75", high = "red") +
    geom_hline(yintercept = 0.6, color = "yellow") +
    geom_abline(slope = 0.001, intercept = 0.1, color = "blue", lty = "93133313")
```




```
## Saving 6.5 x 4.5 in image
```

1.3 Many plotting options

Statistics can be added to the plot as an additional layer. Other layers are coordinates, facets, and scales.

```
ggplot(data = boston) +
  geom_point(mapping = aes(x = rm, y = medv, color = crim), alpha=0.75) +
  geom_smooth(mapping = aes(x = rm, y = medv)) +
  coord_cartesian(xlim = c(4.5, 7.5)) +
  scale_y_log10() +
  scale_color_gradient(low = "yellow", high = "blue") +
  labs(x = "average rooms / house", y = "median house price ($K)",
       title = "Boston median house prices vs. average house size")
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Maybe more useful if colored by quantile of `crim` value.

```
boston %>%
  mutate(crim = cume_dist(crim)) %>%
  ggplot() +
    geom_point(mapping = aes(x = rm, y = medv, color = crim), alpha=0.75) +
    geom_smooth(mapping = aes(x = rm, y = medv)) +
    coord_cartesian(xlim = c(4.5, 7.5)) +
    scale_y_log10() +
    scale_color_gradient(low = "yellow", high = "blue") +
    labs(x = "average rooms / house", y = "median house price ($K)",
         title = "Boston median house prices, house size, and crime")
```

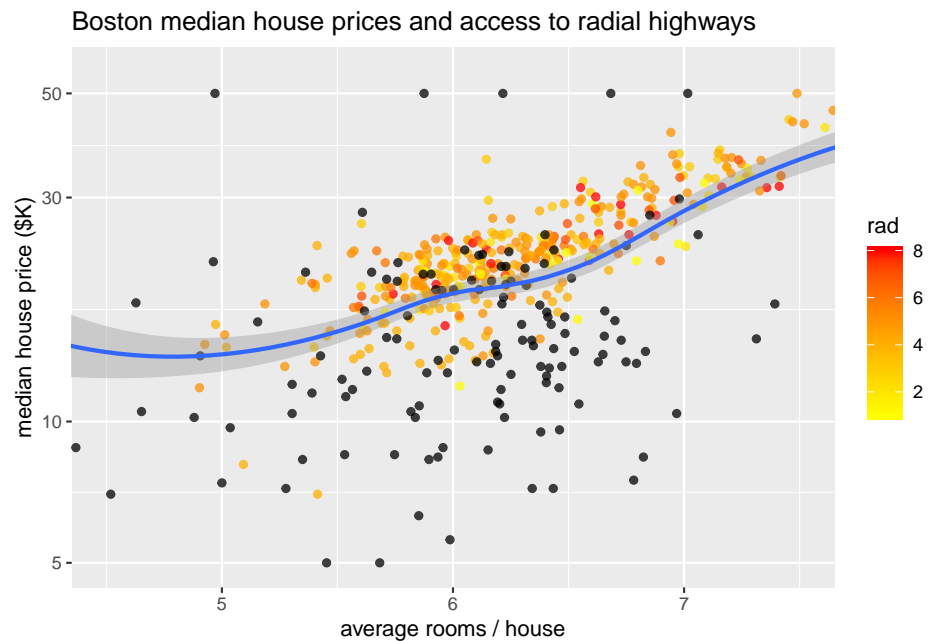
```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Now color by `rad` but change all 24's to NA's.

```
boston %>%
  mutate(rad = ifelse(rad == 24, NA, rad)) %>%
  ggplot() +
    geom_point(mapping = aes(x = rm, y = medv, color = rad), alpha=0.75) +
    geom_smooth(mapping = aes(x = rm, y = medv)) +
    coord_cartesian(xlim = c(4.5, 7.5)) +
    scale_y_log10() +
    scale_color_gradient(low = "yellow", high = "red", na.value = "black") +
    labs(x = "average rooms / house", y = "median house price ($K)",
         title = "Boston median house prices and access to radial highways")
```

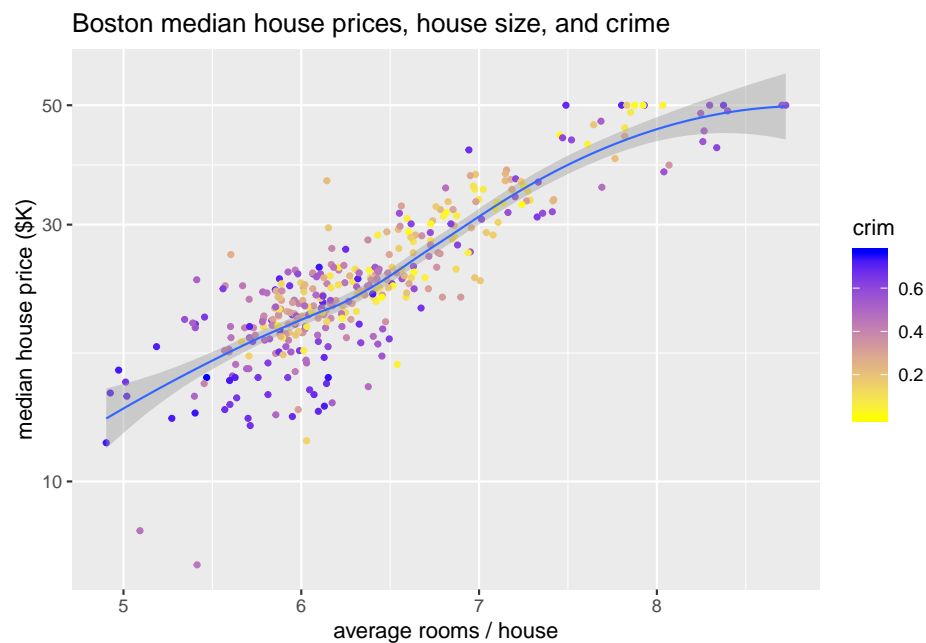
```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Maybe excluding newly-NA'ed rad values helps the crime plot.

```
boston %>%
  mutate(rad = ifelse(rad == 24, NA, rad)) %>%
  mutate(crim = cume_dist(crim)) %>%
  filter(!is.na(rad)) %>%
  ggplot() +
    geom_point(mapping = aes(x = rm, y = medv, color = crim), size = 1) +
    geom_smooth(mapping = aes(x = rm, y = medv), lwd = 0.5) +
    scale_y_log10() +
    scale_color_gradient(low = "yellow", high = "blue") +
    labs(x = "average rooms / house", y = "median house price ($K)",
         title = "Boston median house prices, house size, and crime")
```

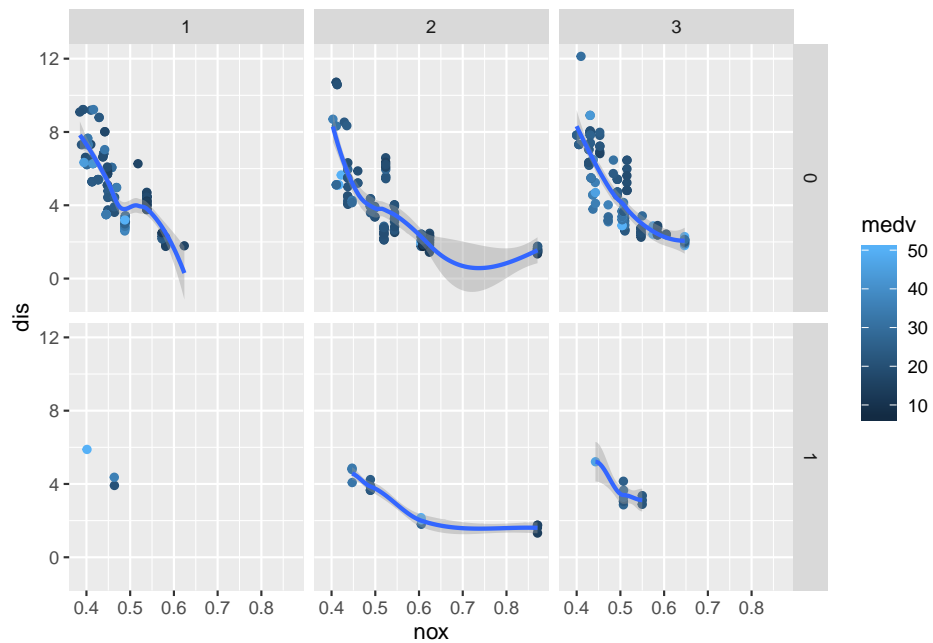
```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

A grid of `nox` vs. `dis` plots according to `chas` (rows) and binned level (`ntile`) of `rad`.


```
boston %>%
  mutate(rad = ifelse(rad == 24, NA, rad)) %>%
  filter(!is.na(rad)) %>%
  ggplot(aes(nox, dis, color = medv)) +
    geom_jitter() +
    facet_grid(chas ~ ntile(rad, 3)) +
    geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




Multiplots available with `gridExtra`, used by `ggplot2`.

```
require(gridExtra)
p1 <- ggplot(boston) +
  geom_point(aes(nox, medv, shape = chas, alpha = cume_dist(lstat)),
    color = 'red', size = 2) +
  labs(title = 'using aes(alpha, shape)')
p2 <- boston %>%
  mutate(lstat_gr = ntile(lstat, 3)) %>%
  ggplot(aes(tax, medv, color = lstat_gr, size = nox)) +
  geom_point(shape = 16, alpha = 0.75) +
  geom_smooth(aes(group = lstat_gr), lwd = 0.8) +
  labs(title = 'using aes(group, size)')
p1
```



p2

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggsave('plots/two-plot.pdf', arrangeGrob(p1, p2))

## Saving 6.5 x 4.5 in image
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

Chapter 2

dplyr and tidyr

```
library(tidyverse)
library(gridExtra)
batting <- as_tibble(Lahman::Batting)
fielding <- as_tibble(Lahman::Fielding)
```

2.1 Hoofin' it with dplyr

Condense batting stats into player career totals, keep only those ≥ 500 games.

```
is_col <- names(batting)[c(1, 2, 4, 6:17)]
is_num <- names(batting)[sapply(batting, is.numeric)]
gt_500 <- batting %>%
  select(is_col) %>%
  select(-teamID) %>%
  drop_na() %>%
  group_by(playerID) %>%
  summarize_at(is_col[-(1:3)], sum, na.rm = T) %>%
  filter(G >= 500)
```

All Ha~ Green~ statistics:

```
batting %>%
  filter(str_detect(playerID, "greenha"))
```

```
## # A tibble: 14 x 22
##   playerID yearID stint teamID lgID      G    AB    R    H   X2B   X3B
##   <chr>      <int> <int> <fct> <fct> <int> <int> <int> <int> <int> <int>
## 1 greenha~  1930     1 DET    AL      1     1     0     0     0     0
## 2 greenha~  1933     1 DET    AL    117   449    59   135    33     3
```

```
## 3 greenha~ 1934 1 DET AL 153 593 118 201 63 7
## 4 greenha~ 1935 1 DET AL 152 619 121 203 46 16
## 5 greenha~ 1935 1 BRO NL 2 0 0 0 0 0
## 6 greenha~ 1936 1 DET AL 12 46 10 16 6 2
## 7 greenha~ 1937 1 DET AL 154 594 137 200 49 14
## 8 greenha~ 1938 1 DET AL 155 556 144 175 23 4
## 9 greenha~ 1939 1 DET AL 138 500 112 156 42 7
## 10 greenha~ 1940 1 DET AL 148 573 129 195 50 8
## 11 greenha~ 1941 1 DET AL 19 67 12 18 5 1
## 12 greenha~ 1945 1 DET AL 78 270 47 84 20 2
## 13 greenha~ 1946 1 DET AL 142 523 91 145 29 5
## 14 greenha~ 1947 1 PIT NL 125 402 71 100 13 2
## # ... with 11 more variables: HR <int>, RBI <int>, SB <int>, CS <int>,
## # BB <int>, SO <int>, IBB <int>, HBP <int>, SH <int>, SF <int>,
## # GIDP <int>
```

Positions by game.

```
fielding %>%
  group_by(POS) %>%
  count(wt = G)
```

```
## # A tibble: 7 x 2
## # Groups:   POS [7]
##   POS      n
##   <chr> <int>
## 1 1B    482698
## 2 2B    480968
## 3 3B    482320
## 4 C     497547
## 5 OF    1451301
## 6 P     1106574
## 7 SS     479045
```

Attach a column denoting their main fielding position.

```
is_field = names(fielding)[c(1, 6, 7, 9, 10, 11, 12, 13)]
fielding %>%
  select(is_field) %>%
  map(~ sum(is.na(.)))
```

```
## $playerID
## [1] 0
##
## $POS
## [1] 0
##
## $G
```

```
## [1] 0
##
## $InnOuts
## [1] 29929
##
## $P0
## [1] 0
##
## $A
## [1] 0
##
## $E
## [1] 1
##
## $DP
## [1] 0
```

Removing InnOuts is a good idea, too many missing.

```
is_field = names(fielding)[c(1, 6, 7, 10, 11, 12, 13)]
pos_tot <- fielding %>%
  select(is_field) %>%
  drop_na() %>%
  group_by(playerID, POS) %>%
  summarize_all(sum) %>%
  ungroup() %>%
  filter(G >= 100) %>%
  arrange(playerID, desc(G)) %>%
  group_by(playerID) %>%
  mutate(pos1 = first(POS)) %>%
  filter(POS == pos1) %>%
  select(-pos1)
```

2.2 tidy and relational data

Add fielding info to batting tibble.

```
(batpos <- gt_500 %>%
  left_join(pos_tot, by = "playerID"))
```

```
## # A tibble: 2,667 x 19
##   playerID  G.x  AB    R    H  X2B  X3B   HR  RBI   SB   CS
##   <chr>    <int> <int> <int> <int> <int> <int> <int> <int> <int> <int>
## 1 aaronha~  3298 12364 2174 3771  624   98  755 2297  240  73
## 2 abbotku~   702  2044  273  523  109   23   62  242   22  11
```

```
## 3 abernte~ 681 181 12 25 3 0 0 9 0 0
## 4 abramca~ 521 1543 246 422 62 19 32 134 11 18
## 5 abreubo~ 2425 8480 1453 2470 574 59 288 1363 400 128
## 6 abreujo~ 742 2913 398 858 180 13 146 488 8 3
## 7 ackledu~ 635 2125 261 512 94 18 46 216 31 12
## 8 adairje~ 1165 4019 378 1022 163 19 57 366 29 29
## 9 adamsbo~ 797 2604 395 701 107 31 25 188 25 30
## 10 adamsgl~ 661 1617 152 452 79 5 34 225 6 10
## # ... with 2,657 more rows, and 8 more variables: BB <int>, SO <int>,
## # POS <chr>, G.y <int>, PO <int>, A <int>, E <int>, DP <int>
```

Counts of positions.

```
batpos %>%
  group_by(POS) %>%
  count()
```

```
## # A tibble: 8 x 2
## # Groups:   POS [8]
##   POS      n
##   <chr> <int>
## 1 <NA>     2
## 2 1B      254
## 3 2B      277
## 4 3B      270
## 5 C       300
## 6 OF      890
## 7 P       378
## 8 SS      296
```

NAs are likely DHs.

```
pos_nas <- batpos %>%
  filter(is.na(POS))
batting %>%
  inner_join(pos_nas, by = "playerID")
```

```
## # A tibble: 26 x 40
##   playerID yearID stint teamID lgID      G AB.x  R.x  H.x X2B.x X3B.x
##   <chr>      <int> <int> <fct> <fct> <int> <int> <int> <int> <int> <int>
## 1 moraljo~ 1973     1 OAK  AL     6  14    0    4    1    0
## 2 moraljo~ 1973     2 MON  NL     5   5    0    2    0    0
## 3 moraljo~ 1974     1 MON  NL    25  26    3    7    4    0
## 4 moraljo~ 1975     1 MON  NL   93 163   18   49    6    1
## 5 moraljo~ 1976     1 MON  NL  104 158   12   50   11    0
## 6 moraljo~ 1977     1 MON  NL   65  74    3   15    4    1
## 7 moraljo~ 1978     1 MIN  AL  101 242   22   76   13    1
## 8 moraljo~ 1979     1 MIN  AL   92 191   21   51    5    1
```



```
## 9 moraljo~ 1980      1 MIN    AL      97   241   36   73   17   2
## 10 moraljo~ 1981      1 BAL    AL      38    86    6   21    3   0
## # ... with 16 more rows, and 29 more variables: HR.x <int>, RBI.x <int>,
## #   SB.x <int>, CS.x <int>, BB.x <int>, SO.x <int>, IBB <int>, HBP <int>,
## #   SH <int>, SF <int>, GIDP <int>, G.x <int>, AB.y <int>, R.y <int>,
## #   H.y <int>, X2B.y <int>, X3B.y <int>, HR.y <int>, RBI.y <int>,
## #   SB.y <int>, CS.y <int>, BB.y <int>, SO.y <int>, POS <chr>, G.y <int>,
## #   PO <int>, A <int>, E <int>, DP <int>
```


Drop these two DHs.

```
batpos <- batpos %>%
  drop_na()
```

Now we could explore many aspects of hitting stats vs. position, and see what position players were better fielders or better hitters, or if neither we can see if they played for the Expos.


```
batpos %>%
  filter(POS == "SS") %>%
  mutate(BA = H / AB) %>% # batting average, hits / at bats
  mutate(Err = E / (PO + A)) %>% # error rate, errors / (put outs + assists)
  mutate(HRR = HR / AB) %>% # home run rate, home runs / at bats
  ggplot(aes(Err, BA)) +
    geom_point(aes(color = HRR)) +
    geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
temp <- batpos %>%
  mutate(BA = H / AB) %>% # batting average, hits / at bats
  filter(between(BA, 0.01, 0.49)) %>%
  mutate(Err = E / (PO + A)) %>% # error rate, errors / (put outs + assists)
  mutate(HRR = HR / AB) # home run rate, home runs / at bats
p1 <- temp %>%
  ggplot(aes(Err, BA, color = POS)) +
  geom_point(alpha = 0.5, size = 0.5) +
  geom_smooth(aes(group = POS)) +
  coord_cartesian(xlim = c(0, 0.1), ylim = c(0.1, 0.42))
p2 <- temp %>%
  filter(POS != "p") %>%
  ggplot(aes(BA, HRR, color = POS)) +
  geom_point(alpha = 0.5, size = 0.5) +
  geom_smooth(aes(group = POS))
p1
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



p2

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggsave('plots/pos-bat.pdf', arrangeGrob(p1, p2))

## Saving 6.5 x 4.5 in image
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

Bibliography

Xie, Y. (2015). *Dynamic Documents with R and knitr*. Chapman and Hall/CRC, Boca Raton, Florida, 2nd edition. ISBN 978-1498716963.

Xie, Y. (2019). *bookdown: Authoring Books and Technical Documents with R Markdown*. R package version 0.11.